

Table of Contents

1	AI-Augmented Development Methodology	1
1.1	Introduction: State of the Art Meets Production Reality	1
1.2	Part 1: Depth of Expertise vs. Internet Noise	2
1.2.1	Investment Profile	2
1.2.2	Claude Code CLI Specialization	2
1.2.3	AI Prompt Engineering, Context Coding, and “AI Vibing”	2
1.3	Part 2: Strategic Framework—Why Software Robustness Requires Different Principles than Business Advantage	3
1.3.1	Classical Strategy vs. Modern Software Engineering	3
1.3.2	The Three-Axis Framework for AI-Augmented Development	3
1.3.3	Design Rule for Effective AI Agent Usage	4
1.3.4	Implications for Netstrata’s Software Development	4
1.4	Part 3: Knowledge Transfer and Team Capability Building	5
1.4.1	Training Offerings for Netstrata’s Software Team	5
1.4.2	Realistic Expectations	5
1.5	Part 4: Assessment and Adoption Strategy	6
1.5.1	Acknowledging Industry Skepticism	6
1.5.2	Demonstration-First Approach	6
1.5.3	Why This Matters for Netstrata’s Timing	6
1.5.4	Mutual Assessment Process	7
1.5.5	No Pressure, Just Results	7
1.5.6	Team-Wide Effectiveness vs. Individual Heroics	8
1.6	Conclusion: Methodology Informs Contribution	8

1 AI-Augmented Development Methodology

1.1 Introduction: State of the Art Meets Production Reality

Modern software development has reached an inflection point. AI coding agents—particularly Anthropic’s Claude Code CLI and similar tools—have transitioned from experimental curiosities to production-grade development accelerators. However, the gap between casual experimentation and effective production usage is measured in thousands of hours of practice, not blog posts written after weekend trials.

This document outlines my approach to AI-augmented development workflows, the economic rationale behind systematic tooling adoption, and the professional requirements necessary for these methodologies to succeed in team environments.

1.2 Part 1: Depth of Expertise vs. Internet Noise

1.2.1 Investment Profile

Thousands of hours of hands-on practice across multiple AI coding platforms:

- **Claude Code CLI:** Extensive Agent Skills development, custom workflow design, and adoption advocacy
- **Production deployment:** Not prototyping or experimentation—these tools run critical business workflows daily
- **Cross-platform experience:** Claude Code, GitHub Copilot, Cursor, and other AI-augmented IDEs
- **Production systems built:** Multiple business-critical systems developed entirely using AI-augmented workflows

Contrast with industry discourse: Most commentary about AI coding agents comes from developers who've used them casually for days or weeks. The noise-to-signal ratio is extremely high—many strong opinions based on minimal practice. Thousands of hours reveal capabilities and limitations that aren't visible in surface-level usage.

1.2.2 Claude Code CLI Specialization

Agent Skills Development: Active contribution to Claude Code ecosystem through:

- Custom Agent Skills creation for specialized workflows
- Internal popularization at Eon Labs (team adoption and training)
- Deep understanding of Claude Code's tool use patterns, context window management, and autonomous agent capabilities
- Experience with both successful and failed automation attempts—knowing when AI tools add value vs. when they introduce unnecessary complexity

What “mastery” actually means:

- Understanding prompt engineering patterns that produce consistent results
- Recognizing when to use AI-assisted coding vs. when manual implementation is faster
- Designing context-aware workflows that leverage large context windows effectively
- Building reliable systems on top of non-deterministic AI outputs
- Managing API costs through efficient prompt design and caching strategies

1.2.3 AI Prompt Engineering, Context Coding, and “AI Vibing”

Prompt Engineering: Systematic approaches to eliciting specific behaviors:

- Task decomposition for complex multi-step operations
- Constraint specification (style guides, architectural patterns, testing requirements)
- Examples-based learning (few-shot prompting with representative cases)
- Iterative refinement based on output quality metrics

Context Coding: Leveraging large context windows (200,000+ tokens):

- Feeding entire codebases into context for architectural consistency
- Processing full specification documents during implementation
- Analyzing large log files or datasets for debugging
- Maintaining conversation state across extended development sessions

AI Vibing: Collaborative development patterns where AI suggestions guide exploration:

- Rapid prototyping through conversational iteration
- Exploration of alternative implementations
- Real-time code review and refactoring suggestions
- Documentation generation that reflects actual implementation

Critical insight: These aren't separate techniques—they're integrated practices that compound over time. The developer who's spent thousands of hours develops intuition about which patterns work in which contexts.

1.3 Part 2: Strategic Framework—Why Software Robustness Requires Different Principles than Business Advantage

1.3.1 Classical Strategy vs. Modern Software Engineering

In classical strategy, abnormal profits arise when a firm controls resources and capabilities that are valuable, rare, and difficult to imitate. These often take the form of proprietary knowledge, data, and processes that competitors cannot easily access or replicate.

By contrast, in modern software engineering practice, robustness and maintainability often come from conformity to shared idioms: widely used languages, standard libraries, and community-validated patterns.

The fundamental tension: In most businesses, margins come from secrets—proprietary knowledge, non-obvious insights, and idiosyncratic processes create advantage. That is **not** how modern software engineering works.

In software today, the opposite is usually true: the more we conform to shared idioms, the **more future-proof** we become.

- Idiomatic UI patterns → easier for users to understand and designers to extend
- Idiomatic language features → easier for other engineers and tools to reason about
- Widely adopted libraries used idiomatically → easier to hire for, easier to maintain, easier to replace

So while business advantage is often *bespoke*, software robustness and maintainability are often *idiomatic*.

1.3.2 The Three-Axis Framework for AI-Augmented Development

Through thousands of hours of production AI agent usage, I've identified three critical tensions not adequately addressed in current industry discourse.

We write **A → B** to mean: “A is constrained / bounded by B.”

Exploration → Specification: The agent’s ability to explore new methods, libraries, and architectures is constrained by how prescriptive we are in the specification. If we stuff the prompt with bespoke, highly specific implementation details, we increase Specification and shrink Exploration. We prevent the agent from roaming the current ecosystem, discovering state-of-the-art, future-proof idioms.

Autonomy → Supervision: The agent’s autonomy is constrained by the level and frequency of human supervision. If we intervene at every small decision, require approvals at every step, or constantly redirect it, we reduce Autonomy. That defeats the point of using an autonomous AI to free up human time.

Idiomatic patterns → Bespoke constraints: The agent’s ability to choose idiomatic, community-tested patterns is constrained by bespoke constraints we impose. If we insist on custom frameworks, unusual architectures, and local “house styles” that diverge from the ecosystem, we block it from using the idioms that maximize compatibility, tooling support, and long-term maintainability.

1.3.3 Design Rule for Effective AI Agent Usage

If we overload an AI agent with bespoke instructions and micro-specifications, we do two things at once:

1. We import the “business secret sauce” mindset into a domain where it is often harmful—telling the agent to ignore the ecosystem and instead re-implement our quirks.
2. We mathematically reduce the left-hand side of each relation: heavy specification throttles exploration, frequent supervision throttles autonomy, bespoke constraints throttle idiomaticity.

As AI coding agents get smarter, the optimal balance tilts further left. We as humans must learn to let go while keeping just the right elements on the right side.

The design rule: - Keep goals, constraints, and safety requirements **clear** - Keep implementation details **non-bespoke and non-prescriptive** where possible

1.3.4 Implications for Netstrata’s Software Development

For software engineering agents working toward the end-2026 completion deadline:

Maximize Exploration: Let agents discover current best practices for NSW compliance automation, testing frameworks, and deployment patterns rather than prescribing outdated approaches.

Maximize Autonomy: Reduce intervention overhead—trust agents to make routine implementation decisions, reserving human judgment for architectural choices and domain-specific NSW regulatory requirements.

Maximize Idiomatic Patterns: Competitive advantage comes from domain knowledge (decade of encoded NSW regulatory expertise), not from custom frameworks. Use community-standard tools idiomatically to maximize maintainability of the \$12-14M investment.

The insight: Netstrata's moat is the NSW strata management knowledge embedded in the software, not the technical implementation patterns. Letting AI agents apply idiomatic, state-of-the-art software engineering practices protects that investment better than insisting on bespoke technical approaches.

1.4 Part 3: Knowledge Transfer and Team Capability Building

1.4.1 Training Offerings for Netstrata's Software Team

Onboarding Workshops: Structured introduction to AI-augmented workflows:

- Tool installation and configuration (Claude Code CLI setup)
- Basic prompt engineering patterns for common development tasks
- Context window management strategies
- Integration with existing development workflows

Pair Programming Sessions: Hands-on learning through collaboration:

- Real-world problem solving using AI-augmented approaches
- Demonstrating when AI tools add value vs. when manual work is faster
- Building intuition through repetitive practice
- Troubleshooting common failure modes

Custom Workflow Design: Tailoring AI tools to Netstrata's specific needs:

- Identifying high-value automation opportunities in current development process
- Building custom Agent Skills for recurring tasks
- Establishing team conventions for AI tool usage
- Creating feedback loops for continuous improvement

Assessment and Iteration: Not one-time training, but ongoing capability building:

- Measuring adoption rates and productivity impacts
- Identifying barriers to effective usage
- Adjusting workflows based on team feedback
- Building internal champions for best practices

1.4.2 Realistic Expectations

Not a silver bullet: AI coding agents don't replace software engineering judgment. They augment it.

Learning curve exists: Team members need weeks of practice to become proficient, months to develop mastery.

Cultural fit matters: Some developers embrace these tools quickly; others resist. Effective adoption requires collective buy-in.

1.5 Part 4: Assessment and Adoption Strategy

1.5.1 Acknowledging Industry Skepticism

I understand the common view that AI coding tools are proof-of-concept only. Most industry commentary—blog posts, YouTube reviews, conference talks—shares this skepticism. Many practitioners have experimented casually and concluded these tools aren't production-ready.

This skepticism is reasonable given surface-level usage. Developers who spend a weekend with Copilot or try Claude Code for a few hours often encounter limitations that reinforce doubts about production readiness. The tools CAN be frustrating when used naively.

However, thousands of hours of deep practice reveal different patterns: these tools ARE production-ready when integrated systematically into team workflows. The gap between “tried it casually” and “production mastery” is substantial—similar to the difference between writing your first Python script and architecting production systems.

1.5.2 Demonstration-First Approach

Rather than advocate for adoption based on theory or external claims, I propose demonstration-based assessment for Netstrata:

First 30-60 Days: Work with Tom Bakani's software team on Phase 1 completion tasks using AI-augmented workflows. Approach can be tailored based on team's current priorities and the specific challenges they're facing.

Team Observation: Pair programming sessions and live demonstrations where team members can see these tools in action on the actual Netstrata codebase. Real code, real challenges, real results—not theoretical examples.

Measurable Outcomes: Document time savings, code quality improvements, and velocity gains compared to traditional development methods. Quantifiable data rather than subjective impressions.

Results-Based Decision: After seeing production usage firsthand over multiple weeks, the team assesses whether these approaches deliver value for Netstrata's priorities. If yes, we explore broader adoption pathways. If skepticism remains after demonstrations, we have objective data to inform decisions.

1.5.3 Why This Matters for Netstrata's Timing

Competitive Window: Strata software competitors (Strata Master, Different, external vendors) are unlikely to be using cutting-edge AI development practices yet. Industry-wide skepticism means early adopters gain 12-24 months of velocity advantage.

Critical Insight: By the time industry consensus shifts to “these tools are production-ready” (which may take 12-24 months given current skepticism levels), Netstrata will already have that time period’s worth of accumulated advantage. You’re not adopting unproven technology—you’re adopting ahead of the curve while competitors wait for consensus.

End-2026 Deadline Pressure: Phase 1 completion timeline benefits immediately from accelerated development velocity. This isn’t theoretical future benefit—it’s practical advantage on current priorities with concrete deadlines.

Protection of \$12-14M Investment: AI-assisted code review, comprehensive testing, and systematic quality assurance become especially valuable as the codebase approaches completion. The cost of bugs and technical debt increases as you get closer to external rollout.

1.5.4 Mutual Assessment Process

This opportunity requires fit assessment on both sides:

For Netstrata to Assess:

- Do AI-augmented workflows deliver measurable value for software completion timeline?
- Does the velocity improvement justify the learning investment for Tom Bakani’s team?
- Are the competitive timing advantages real for the strata software market?

For Me to Assess:

- Is the team open to exploring modern development practices through demonstration?
- Will leadership (Tom Bakani, Andrew Tunks) support capability-building efforts?
- Can we establish feedback loops for continuous improvement?

Better to identify misalignment through demonstration and data than assume compatibility or incompatibility without evidence.

1.5.5 No Pressure, Just Results

I’m not advocating that everyone immediately become an AI coding expert or overhaul their workflows on day one. I’m offering to:

1. **Demonstrate capabilities** on real Netstrata work (Phase 1 completion tasks)
2. **Document results** objectively (time measurements, quality metrics, velocity data)
3. **Let the team decide** based on evidence rather than theory

If **Tom Bakani’s team sees value** after demonstrations and wants to explore adoption, I’m equipped to support that transition through the training and knowledge transfer approaches outlined in Part 3.

If **they remain unconvinced** after seeing production usage over 30-60 days, that’s valuable signal about fit. No hard feelings—some teams prefer traditional workflows, and that’s a legitimate choice when made with full information.

1.5.6 Team-Wide Effectiveness vs. Individual Heroics

One observation from extensive experience: AI-augmented development works best as a team practice rather than individual capability. When one person uses these tools and others don't, it can create coordination friction (different development speeds, code review challenges, knowledge silos).

However, this doesn't mean "all or nothing" adoption. It DOES mean:

Gradual, organic adoption: Start with demonstrations, let interested team members experiment, build capability progressively. Not forcing anyone, but creating environment where exploration is encouraged.

Leadership support: Tom Bakani and technical leadership actively encouraging the team to evaluate these approaches fairly, not dismissing them as hype.

Patience for learning curves: Understanding that proficiency takes weeks/months of practice, not days. Allow time for team members to develop intuition.

Assessment based on results: If demonstrations show clear value, team explores further. If results are marginal, we reassess. Data-driven decisions.

1.6 Conclusion: Methodology Informs Contribution

This methodology directly shapes how I'd contribute to Netstrata's three phases:

Phase 1 (Software Completion): Demonstrate AI-augmented workflows on actual development tasks, document velocity gains, train interested team members, accelerate progress toward end-2026 deadline.

Phase 2 (WA Migration): Apply AI-assisted documentation generation, validation script creation, and migration infrastructure development to prepare for customer rollout—areas where these tools excel.

Phase 3 (External Rollout): Apply lessons from Phase 1-2 demonstrations to scale modern development practices if the team decides broader adoption delivers value.

The core insight: Technology adoption should be driven by demonstrated results, not theory or hype. For a \$12-14M software investment with strategic importance to Netstrata's competitive position, exploring state-of-the-art development tools through systematic demonstration is a low-risk, high-potential-upside opportunity.

The window for gaining competitive advantage through early adoption is now—while competitors remain skeptical. Let results speak.